

# Muhammad\_Dastgir\_Breast\_Cancer

November 11, 2024

```
[2]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from ucimlrepo import fetch_ucirepo
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LogisticRegressionCV
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import random
```

## 0.1 Loading Dataset from UC Irvine Machine Learning Repository and Assigning Values

```
[5]: breast_cancer_wisconsin_diagnostic = fetch_ucirepo(id=17)
X = breast_cancer_wisconsin_diagnostic.data.features
y = breast_cancer_wisconsin_diagnostic.data.targets
```

```
[6]: X
```

```
[6]:
```

	radius1	texture1	perimeter1	area1	smoothness1	compactness1	\
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	
..	...	...	...	...	...	...	
564	21.56	22.39	142.00	1479.0	0.11100	0.11590	
565	20.13	28.25	131.20	1261.0	0.09780	0.10340	
566	16.60	28.08	108.30	858.1	0.08455	0.10230	
567	20.60	29.33	140.10	1265.0	0.11780	0.27700	
568	7.76	24.54	47.92	181.0	0.05263	0.04362	

  

	concavity1	concave_points1	symmetry1	fractal_dimension1	...	radius3	\
0	0.30010	0.14710	0.2419	0.07871	...	25.380	
1	0.08690	0.07017	0.1812	0.05667	...	24.990	
2	0.19740	0.12790	0.2069	0.05999	...	23.570	
3	0.24140	0.10520	0.2597	0.09744	...	14.910	
4	0.19800	0.10430	0.1809	0.05883	...	22.540	

```

..      ...      ...      ...      ...      ...
564      0.24390      0.13890      0.1726      0.05623 ... 25.450
565      0.14400      0.09791      0.1752      0.05533 ... 23.690
566      0.09251      0.05302      0.1590      0.05648 ... 18.980
567      0.35140      0.15200      0.2397      0.07016 ... 25.740
568      0.00000      0.00000      0.1587      0.05884 ... 9.456

```

```

      texture3  perimeter3  area3  smoothness3  compactness3  concavity3  \
0      17.33      184.60  2019.0      0.16220      0.66560      0.7119
1      23.41      158.80  1956.0      0.12380      0.18660      0.2416
2      25.53      152.50  1709.0      0.14440      0.42450      0.4504
3      26.50      98.87   567.7      0.20980      0.86630      0.6869
4      16.67      152.20  1575.0      0.13740      0.20500      0.4000
..      ...      ...      ...      ...      ...
564      26.40      166.10  2027.0      0.14100      0.21130      0.4107
565      38.25      155.00  1731.0      0.11660      0.19220      0.3215
566      34.12      126.70  1124.0      0.11390      0.30940      0.3403
567      39.42      184.60  1821.0      0.16500      0.86810      0.9387
568      30.37      59.16   268.6      0.08996      0.06444      0.0000

```

```

      concave_points3  symmetry3  fractal_dimension3
0      0.2654      0.4601      0.11890
1      0.1860      0.2750      0.08902
2      0.2430      0.3613      0.08758
3      0.2575      0.6638      0.17300
4      0.1625      0.2364      0.07678
..      ...      ...      ...
564      0.2216      0.2060      0.07115
565      0.1628      0.2572      0.06637
566      0.1418      0.2218      0.07820
567      0.2650      0.4087      0.12400
568      0.0000      0.2871      0.07039

```

[569 rows x 30 columns]

[7]: y

```

[7]:      Diagnosis
0      M
1      M
2      M
3      M
4      M
..      ...
564      M
565      M
566      M

```

```
567         M
568         B
```

```
[569 rows x 1 columns]
```

### 0.1.1 Creating training and testing data

```
[8]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
↳ random_state = 69)
```

```
[9]: training_model = LogisticRegression()
training_model.fit(X_train, y_train)
```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/utils/validation.py:1184: DataConversionWarning: A column-
vector y was passed when a 1d array was expected. Please change the shape of y
to (n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[9]: LogisticRegression()
```

```
[10]: predicted_y = training_model.predict(X_test)
predicted_y # for reference
```

```
[10]: array(['M', 'B', 'M', 'B', 'M', 'M', 'B', 'B', 'B', 'M', 'B', 'M', 'B',
'M', 'B', 'B', 'M', 'M', 'M', 'M', 'B', 'B', 'M', 'M', 'M', 'M',
'M', 'B', 'B', 'M', 'M', 'M', 'B', 'B', 'B', 'M', 'B', 'B', 'M',
'M', 'B', 'B', 'B', 'B', 'B', 'M', 'B', 'B', 'M', 'B', 'B', 'M',
'B', 'M', 'M', 'B', 'B', 'M', 'B', 'B', 'B', 'M', 'B', 'B', 'M',
'B', 'M', 'B', 'B', 'M', 'M', 'B', 'B', 'B', 'B', 'B', 'M', 'B',
'M', 'M', 'B', 'M', 'M', 'B', 'B', 'M', 'B', 'B', 'M', 'B', 'M',
'B', 'B', 'B', 'B', 'B', 'M', 'B', 'B', 'B', 'M', 'M', 'B', 'M',
'M', 'M', 'B', 'B', 'B', 'M', 'M', 'B', 'B', 'M'], dtype=object)
```

```
[11]: accuracy = training_model.score(X_test, y_test)
accuracy
```

```
[11]: 0.9298245614035088
```

```
[12]: Cs = np.logspace(-4, 4, 10)
```

```
[13]: l1_model = LogisticRegressionCV(  
    Cs=Cs,  
    cv=5,  
    penalty='l1',  
    solver='liblinear',  
    max_iter=1000,  
    scoring='accuracy',  
    refit=True  
)
```

```
[14]: l1_model.fit(X_train, y_train)
```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-  
packages/sklearn/utils/validation.py:1184: DataConversionWarning: A column-  
vector y was passed when a 1d array was expected. Please change the shape of y  
to (n_samples, ), for example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-  
packages/sklearn/svm/_base.py:1242: ConvergenceWarning: Liblinear failed to  
converge, increase the number of iterations.
```

```
    warnings.warn(  

```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-  
packages/sklearn/svm/_base.py:1242: ConvergenceWarning: Liblinear failed to  
converge, increase the number of iterations.
```

```
    warnings.warn(  

```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-  
packages/sklearn/svm/_base.py:1242: ConvergenceWarning: Liblinear failed to  
converge, increase the number of iterations.
```

```
    warnings.warn(  

```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-  
packages/sklearn/svm/_base.py:1242: ConvergenceWarning: Liblinear failed to  
converge, increase the number of iterations.
```

```
    warnings.warn(  

```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-  
packages/sklearn/svm/_base.py:1242: ConvergenceWarning: Liblinear failed to  
converge, increase the number of iterations.
```

```
    warnings.warn(  

```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-  
packages/sklearn/svm/_base.py:1242: ConvergenceWarning: Liblinear failed to  
converge, increase the number of iterations.
```

```
    warnings.warn(  

```

```
[14]: LogisticRegressionCV(Cs=array([1.00000000e-04, 7.74263683e-04, 5.99484250e-03,  
    4.64158883e-02,
```

```
3.59381366e-01, 2.78255940e+00, 2.15443469e+01, 1.66810054e+02,
1.29154967e+03, 1.00000000e+04]),
      cv=5, max_iter=1000, penalty='l1', scoring='accuracy',
      solver='liblinear')
```

```
[15]: l2_model = LogisticRegressionCV(
      Cs=Cs,
      cv=5,
      penalty='l2',
      solver='lbfgs',
      max_iter=1000,
      scoring='accuracy',
      refit=True
    )
```

```
[16]: l2_model.fit(X_train, y_train)
```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/utils/validation.py:1184: DataConversionWarning: A column-
vector y was passed when a 1d array was expected. Please change the shape of y
to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
  n_iter_i = _check_optimize_result(
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
  n_iter_i = _check_optimize_result(
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[16]: LogisticRegressionCV(Cs=array([1.00000000e-04, 7.74263683e-04, 5.99484250e-03,
4.64158883e-02,
3.59381366e-01, 2.78255940e+00, 2.15443469e+01, 1.66810054e+02,
1.29154967e+03, 1.00000000e+04]),
cv=5, max_iter=1000, scoring='accuracy')
```

```
[17]: l2_model.fit(X_train, y_train)
```

```
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/utils/validation.py:1184: DataConversionWarning: A column-
vector y was passed when a 1d array was expected. Please change the shape of y
to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/Users/hamiddastgir/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
regression
```

```
    n_iter_i = _check_optimize_result(
```

```
[17]: LogisticRegressionCV(Cs=array([1.00000000e-04, 7.74263683e-04, 5.99484250e-03,
4.64158883e-02,
      3.59381366e-01, 2.78255940e+00, 2.15443469e+01, 1.66810054e+02,
      1.29154967e+03, 1.00000000e+04]),
      cv=5, max_iter=1000, scoring='accuracy')
```

```
[18]: optimal_C_l1 = l1_model.C_[0]
      optimal_C_l2 = l2_model.C_[0]

      print(f'Optimal C for L1 penalty: {optimal_C_l1}')
      print(f'Optimal C for L2 penalty: {optimal_C_l2}')
```

```
Optimal C for L1 penalty: 21.54434690031882
```

```
Optimal C for L2 penalty: 1291.5496650148827
```

```
[19]: y_pred_l1 = l1_model.predict(X_test)
      y_pred_l2 = l2_model.predict(X_test)

      accuracy_l1 = accuracy_score(y_test, y_pred_l1)
      accuracy_l2 = accuracy_score(y_test, y_pred_l2)

      print(f'L1 Penalty Accuracy: {accuracy_l1:.2f}')
      print(f'L2 Penalty Accuracy: {accuracy_l2:.2f}')
```

```
L1 Penalty Accuracy: 0.95
```

```
L2 Penalty Accuracy: 0.93
```

```
[20]: n_features_l1 = np.sum(l1_model.coef_ != 0)
      n_features_l2 = np.sum(l2_model.coef_ != 0)

      print(f'Number of features used in L1 model: {n_features_l1}')
      print(f'Number of features used in L2 model: {n_features_l2}')
```

```
Number of features used in L1 model: 18
```

```
Number of features used in L2 model: 30
```

## 0.2 Analysis

The lasso regularized model was more accurate, and it had more non-zero coefficients. The lasso regularized model also benefitted from a higher lamda (inverse C - used above)

```
[21]: y_pred_l1
```

```
[21]: array(['M', 'M', 'M', 'B', 'M', 'M', 'B', 'B', 'B', 'M', 'B', 'M', 'B',
      'M', 'B', 'M', 'B', 'M', 'M', 'M', 'B', 'B', 'M', 'M', 'M', 'M',
      'M', 'B', 'B', 'M', 'M', 'M', 'B', 'B', 'B', 'M', 'B', 'B', 'M',
      'M', 'B', 'B', 'B', 'B', 'B', 'M', 'M', 'B', 'M', 'B', 'B', 'M',
```



```
'B', 'M', 'M', 'B', 'B', 'M', 'B', 'B', 'B', 'M', 'B', 'B', 'M',
'B', 'M', 'B', 'B', 'M', 'M', 'M', 'B', 'B', 'B', 'B', 'B', 'B',
'M', 'M', 'B', 'M', 'M', 'B', 'B', 'M', 'B', 'B', 'M', 'B', 'M',
'B', 'B', 'B', 'B', 'B', 'B', 'B', 'B', 'M', 'M', 'M', 'B', 'M',
'M', 'M', 'B', 'B', 'B', 'M', 'M', 'B', 'B', 'M'], dtype=object)
```

```
[22]: y_pred_12
```

```
[22]: array(['M', 'M', 'M', 'B', 'M', 'M', 'B', 'B', 'B', 'M', 'B', 'M', 'B',
'M', 'B', 'B', 'B', 'M', 'M', 'M', 'B', 'B', 'M', 'M', 'M', 'M',
'M', 'B', 'B', 'M', 'M', 'M', 'B', 'B', 'B', 'M', 'B', 'B', 'M',
'M', 'B', 'B', 'B', 'B', 'B', 'M', 'M', 'B', 'M', 'B', 'B', 'M',
'B', 'M', 'M', 'B', 'B', 'M', 'B', 'B', 'B', 'M', 'B', 'B', 'M',
'B', 'M', 'B', 'B', 'M', 'M', 'B', 'B', 'B', 'B', 'B', 'M', 'B',
'M', 'M', 'B', 'M', 'M', 'B', 'B', 'M', 'M', 'B', 'M', 'B', 'M',
'B', 'B', 'B', 'B', 'B', 'M', 'B', 'B', 'B', 'M', 'M', 'B', 'M',
'M', 'M', 'B', 'B', 'B', 'M', 'M', 'B', 'B', 'M'], dtype=object)
```